

# DB Final Project

Team 15

104062203 陳涵宇

104062232 廖子毅

104062234 林士軒

# Outline

- Paper introduction
  - Motivation and Problem
  - Main Idea
  - Conclusion
- Implementation
- Evaluation and Experiments
- Conclusion

# Outline

- Paper introduction
  - Motivation and Problem
  - Main Idea
  - Conclusion
- Implementation
- Evaluation and Experiments
- Conclusion

# On Optimistic Methods for Concurrency Control

ACM Transactions on Database System

June 1981

Citation count: 1678

H.T. KUNG and JOHN t. ROBINSON

Carnegie-Mellon University

# Motivation and Problem

- Most current approaches to concurrency control in database systems rely on locking of data objects as a control mechanism.
- Applications for which nonblocking concurrency controls should be more efficient than locking are discussed.
- Assumption: multiple transactions can frequently complete without interfering with each other.

# Outline

- Paper introduction
  - Motivation and Problem
  - Main Idea
  - Conclusion
- Implementation
- Evaluation and Experiments
- Conclusion

# Main Idea

- divided transaction into three phases: read phase, validation phase, write phase




Fig. 1. The three phases of a transaction.

# Main Idea

- read phase: all writes take place on local copies of the nodes to be modified. create and maintain four sets.

```
tcreate = (  
  n := create;  
  create set := create set  $\cup$  {n};  
  return n)  
  
twrite(n, i, v) = (  
  if n  $\in$  create set  
    then write(n, i, v)  
  else if n  $\in$  write set  
    then write(copies[n], i, v)  
  else (  
    m := copy(n);  
    copies[n] := m;  
    write set := write set  $\cup$  {n};  
    write(copies[n], i, v)))
```

```
tread(n, i) = (  
  read set := read set  $\cup$  {n};  
  if n  $\in$  write set  
    then return read(copies[n], i)  
  else  
    return read(n, i))  
  
tdelete(n) = (  
  delete set := delete set  $\cup$  {n}).
```



# Main Idea

- validation phase: the changes of the transaction made will not cause a loss of integrity.

- (1)  $T_i$  completes its write phase before  $T_j$  starts its read phase.
- (2) The write set of  $T_i$  does not intersect the read set of  $T_j$ , and  $T_i$  completes its write phase before  $T_j$  starts its write phase.
- (3) The write set of  $T_i$  does not intersect the read set or the write set of  $T_j$ , and  $T_i$  completes its read phase before  $T_j$  completes its read phase.




Fig. 2. Possible interleaving of two transactions.

# Main Idea

- serial validation

```
tbegin = (  
  create set := empty;  
  read set := empty;  
  write set := empty;  
  delete set := empty;  
  start tn := tnc)
```

# Main Idea

- serial validation

```
tend = (  
  ⟨finish tn := tnc;  
  valid := true;  
  for t from start tn + 1 to finish tn do  
    if (write set of transaction with transaction number t intersects read set)  
      then valid := false;  
  if valid  
    then ((write phase); tnc := tnc + 1; tn := tnc));  
  if valid  
    then (cleanup)  
    else (backup)).
```

# Main Idea

- parallel validation

```
tend = (  
  {finish tn := tnc;  
  finish active := (make a copy of active);  
  active := active  $\cup$  {id of this transaction}};  
  valid := true;  
  for t from start tn + 1 to finish tn do  
    if (write set of transaction with transaction number t intersects read set)  
      then valid := false;  
  for i  $\in$  finish active do  
    if (write set of transaction  $T_i$  intersects read set or write set)  
      then valid := false;  
  if valid  
    then (  
      (write phase);  
      (tnc := tnc + 1;  
      tn := tnc;  
      active := active — {id of this transaction});  
      (cleanup))  
    else (  
      (active := active — {id of transaction});  
      (backup)).  
)
```

# Main Idea

- write phase: the local copies are made global

**for**  $n \in \text{write set}$  **do**  $\text{exchange}(n, \text{copies}[n])$ .  
**for**  $n \in \text{delete set}$  **do**  $\text{delete}(n)$ ;

# Outline

- Paper introduction
  - Motivation and Problem
  - Main Idea
  - Conclusion
- Implementation
- Evaluation and Experiments
- Conclusion

# Conclusion

- pros
  - reach high performance in low conflict rate
  - without maintain lock protocol overhead
- cons
  - the DBMS may has extreme low throughput when high conflict rate

# Outline

- Paper introduction
  - Motivation and Problem
  - Main Idea
  - Conclusion
- **Implementation**
- Evaluation and Experiments
- Conclusion



# pre-work

- create three per-transaction sets ( write, read, delete )

```
public Transaction(TransactionMgr txMgr, TransactionLifecycleListener concurMgr,  
    TransactionLifecycleListener recoveryMgr, TransactionLifecycleListener bufferMgr, boolean readOnly,  
    long txNum, long startTn) {  
    this.concurMgr = (ConcurrencyMgr) concurMgr;  
    this.recoveryMgr = (RecoveryMgr) recoveryMgr;  
    this.bufferMgr = (BufferMgr) bufferMgr;  
    this.txNum = txNum;  
    this.startTn = startTn;  
    this.readOnly = readOnly;  
    this.writeSet = new HashMap<RecordField, Constant>();  
    this.readSet = new HashSet<RecordField>();  
    this.deleteSet = new HashSet<RecordField>();  
    this.cCertified = false;  
    this.rCertified = false;  
}
```

```
public class RecordField {  
    public String tblName;  
    public RecordId rid;  
    public String fldName;  
}
```

# pre-work

- create write sets and tnc in txMgr for success committed tx
- create active set using in validation

```
public class TransactionMgr implements TransactionLifecycleListener {  
    private long nextTxNum = 0;  
    private long tnc = 0;  
    // Optimization: Use separate lock for nextTxNum  
    private Object txNumLock = new Object();  
    private Object tncLock = new Object();  
    private HashMap<Long, HashSet<RecordField>> writeSets;  
    private HashSet<Transaction> activeSets;  
    private Object activeLock = new Object();  
}
```

# Read Phase

- four function map to four action
  - insert() : create + write
  - setVal() : write
  - getVal() : read
  - delete() : delete

# insert()

- create + write

```
tcreate = (  
  n := create;  
  create set := create set  $\cup$  {n};  
  return n)
```

```
twrite(n, i, v) = (  
  if n  $\in$  create set  
    then write(n, i, v)  
  else if n  $\in$  write set  
    then write(copies[n], i, v)  
  else (  
    m := copy(n);  
    copies[n] := m;  
    write set := write set  $\cup$  {n};  
    write(copies[n], i, v)))
```

# setVal()

```
twrite(n, i, v) = (  
  if n  $\in$  create set  
    then write(n, i, v)  
  else if n  $\in$  write set  
    then write(copies[n], i, v)  
  else (  
    m := copy(n);  
    copies[n] := m;  
    write set := write set  $\cup$  {n};  
    write(copies[n], i, v)))
```

## RecordFile

```
public void setVal(String fldName, Constant val) {  
    if (tx.isReadOnly() && !isTempTable())  
        throw new UnsupportedOperationException();  
    Type fldType = ti.schema().type(fldName);  
  
    Constant v = val.castTo(fldType);  
    if (Page.size(v) > Page.maxSize(fldType))  
        throw new SchemaIncompatibleException();  
    if (!tx.rollbackCertified() && !tx.commitCertified() && !isTempTable()) {  
        tx.putVal(ti.tableName(), currentRecordId(), fldName, v);  
    } else {  
        rp.setVal(fldName, v);  
    }  
}
```

**modify in tx write set**

## Transaction

```
public void putVal(String tblName, RecordId rid, String fldName, Constant val) {  
    writeSet.put(new RecordField(tblName, rid, fldName), val);  
}
```

# setVal()

```
twrite(n, i, v) = (  
  if n  $\in$  create set  
    then write(n, i, v)  
  else if n  $\in$  write set  
    then write(copies[n], i, v)  
  else (  
    m := copy(n);  
    copies[n] := m;  
    write set := write set  $\cup$  {n};  
    write(copies[n], i, v)))
```

## RecordFile

```
public void setVal(String fldName, Constant val) {  
    if (tx.isReadOnly() && !isTempTable())  
        throw new UnsupportedOperationException();  
    Type fldType = ti.schema().type(fldName);  
  
    Constant v = val.castTo(fldType);  
    if (Page.size(v) > Page.maxSize(fldType))  
        throw new SchemaIncompatibleException();  
    if (!tx.rollbackCertified() && !tx.commitCertified() && !isTempTable()) {  
        tx.putVal(ti.tableName(), currentRecordId(), fldName, v);  
    } else {  
        rp.setVal(fldName, v);  
    }  
}
```

it's not rollback or commit yet

## Transaction

```
public void putVal(String tblName, RecordId rid, String fldName, Constant val) {  
    writeSet.put(new RecordField(tblName, rid, fldName), val);  
}
```

# getVal()

```
tread(n, i) = (  
  read set := read set  $\cup$  {n};  
  if n  $\in$  write set  
    then return read(copies[n], i)  
  else  
    return read(n, i)
```

## RecordFile

```
public Constant getVal(String fldName) {  
  Constant val;  
  val = tx.getVal(ti.tableName(), currentRecordId(), fldName);  
  if (val != null)  
    return val;  
  return rp.getVal(fldName);  
}
```

check if *n* is in write set

## Transaction

```
public Constant getVal(String tblName, RecordId rid, String fldName) {  
  readSet.add(new RecordField(tblName, rid, fldName));  
  return writeSet.get(new RecordField(tblName, rid, fldName));  
}
```

# getVal()

```
tread(n, i) = (  
  read set := read set  $\cup$  {n};  
  if n  $\in$  write set  
    then return read(copies[n], i)  
  else  
    return read(n, i)
```

## RecordFile

```
public Constant getVal(String fldName) {  
  Constant val;  
  val = tx.getVal(ti.tableName(), currentRecordId(), fldName);  
  if (val != null)  
    return val;  
  return rp.getVal(fldName);  
}
```

## Transaction

```
public Constant getVal(String tblName, RecordId rid, String fldName) {  
  readSet.add(new RecordField(tblName, rid, fldName));  
  return writeSet.get(new RecordField(tblName, rid, fldName));  
}
```

add *n* into read set



# delete()

$tdelete(n) = ($   
 $delete\ set := delete\ set \cup \{n\}).$

## Transaction

```
public void delete(String tblName, RecordId rid) {  
    deleteSet.add(new RecordField(tblName, rid, ""));  
}
```

## RecordFile

```
public void delete() {  
    if (tx.isReadOnly() && !isTempTable())  
        throw new UnsupportedOperationException();  
  
    if(!tx.rollbackCertified() && !tx.commitCertified()) {  
        tx.delete(ti.tableName(), currentRecordId());  
    } else {  
        if (fhp == null)  
            fhp = openHeaderForModification();  
  
        // Log that this logical operation starts  
        RecordId deletedRid = currentRecordId();  
        tx.recoveryMgr().logLogicalStart();  
  
        // Delete the current record  
        rp.delete(fhp.getLastDeletedSlot());  
        fhp.setLastDeletedSlot(currentRecordId());  
  
        // Log that this logical operation ends  
        tx.recoveryMgr().logRecordFileDeletionEnd(ti.tableName(),  
            deletedRid.block().number(), deletedRid.id());  
  
        // Close the header (release the header lock)  
        closeHeader();  
    }  
}
```

**modify in tx delete set**

# delete()

$tdelete(n) = ($   
 $delete\ set := delete\ set \cup \{n\}).$

## Transaction

```
public void delete(String tblName, RecordId rid) {  
    deleteSet.add(new RecordField(tblName, rid, ""));  
}
```

**add  $n$  into delete set**

## RecordFile

```
public void delete() {  
    if (tx.isReadOnly() && !isTempTable())  
        throw new UnsupportedOperationException();  
  
    if (!tx.rollbackCertified() && !tx.commitCertified()) {  
        tx.delete(ti.tableName(), currentRecordId());  
    } else {  
  
        if (fhp == null)  
            fhp = openHeaderForModification();  
  
        // Log that this logical operation starts  
        RecordId deletedRid = currentRecordId();  
        tx.recoveryMgr().logLogicalStart();  
  
        // Delete the current record  
        rp.delete(fhp.getLastDeletedSlot());  
        fhp.setLastDeletedSlot(currentRecordId());  
  
        // Log that this logical operation ends  
        tx.recoveryMgr().logRecordFileDeletionEnd(ti.tableName(),  
            deletedRid.block().number(), deletedRid.id());  
  
        // Close the header (release the header lock)  
        closeHeader();  
    }  
}
```

# Validation Phase

- implement parallel version

```
tend = (  
  {finish tn := tnc;  
  finish active := (make a copy of active);  
  active := active  $\cup$  {id of this transaction}};  
  valid := true;  
  for t from start tn + 1 to finish tn do  
    if (write set of transaction with transaction number t intersects read set)  
      then valid := false;  
  for i  $\in$  finish active do  
    if (write set of transaction  $T_i$  intersects read set or write set)  
      then valid := false;  
  if valid  
    then (  
      (write phase);  
      (tnc := tnc + 1;  
      tn := tnc;  
      active := active — {id of this transaction});  
      (cleanup))  
    else (  
      (active := active — {id of transaction});  
      (backup)).  
)
```

# tbegin

## TransactionMgr

```
private Transaction createTransaction(int isolationLevel, boolean readOnly, long txNum) {  
    tx = new Transaction(this, concurMgr, recoveryMgr, bufferMgr, readOnly, txNum, tnc);  
}
```

```
tbegin = (  
    create set := empty;  
    read set := empty;  
    write set := empty;  
    delete set := empty;  
    start tn := tnc)
```

pass tnc to Transaction constructor

## Transaction

```
public Transaction(TransactionMgr txMgr, TransactionLifecycleListener concurMgr,  
    TransactionLifecycleListener recoveryMgr, TransactionLifecycleListener bufferMgr, boolean readOnly,  
    long txNum, long startTn) {  
    this.concurMgr = (ConcurrencyMgr) concurMgr;  
    this.recoveryMgr = (RecoveryMgr) recoveryMgr;  
    this.bufferMgr = (BufferMgr) bufferMgr;  
    this.txNum = txNum;  
    this.startTn = startTn;  
    this.readOnly = readOnly;  
    this.writeSet = new HashMap<RecordField, Constant>();  
    this.readSet = new HashSet<RecordField>();  
    this.deleteSet = new HashSet<RecordField>();  
    this.cCertified = false;  
    this.rCertified = false;  
}
```

# tend

```
tend = (  
  {finish tn := tnc;  
   finish active := (make a copy of active);  
   active := active  $\cup$  {id of this transaction}};  
  valid := true;
```

```
for t from start tn + 1 to finish tn do  
  if (write set of transaction with transaction number t intersects read set)  
    then valid := false;  
for i  $\in$  finish active do  
  if (write set of transaction  $T_i$  intersects read set or write set)  
    then valid := false;
```

```
if valid  
  then (  
    (write phase);  
    {tnc := tnc + 1;  
     tn := tnc;  
     active := active — {id of this transaction}};  
    (cleanup))  
  else (  
    {active := active — {id of transaction}};  
    {backup}}).
```

Transaction

```
public void commit() {
```

validate before write

```
  if(VanillaDb.txMgr().validate(this)) {  
    commitCertify();  
    commitWorkspace();  
    for (TransactionLifecycleListener l : lifecycleListeners)  
      l.onTxCommit(this);  
  
    if (logger.isLoggable(Level.FINE))  
      logger.fine("transaction " + txNum + " committed");  
  
    VanillaDb.txMgr().finishedCommit(writeSet.keySet());  
    VanillaDb.txMgr().endCommit(this);  
  } else {  
    VanillaDb.txMgr().endCommit(this);  
    throw new ValidAbortException("abort tx." + this.txNum + " for no validation");  
  }  
}
```

```
}
```



# validate()

TransactionMgr

```
public boolean validate(Transaction tx) {
    long start = tx.getStartTn();
    long finish = 0;
    HashSet<Transaction> active;
    synchronized (tncLock) {
        synchronized (activeLock) {
            finish = tnc;
            active = new HashSet<Transaction>(activeSets);
            activeSets.add(tx);
        }
    }
    HashSet<RecordField> readSet = tx.getReadSet();
    HashSet<RecordField> writeSet = tx.getWriteSet();

    // write set of transaction from start tn+1 to finish tn interacts read set
    for(long it=start+1; it<=finish; ++it) {
        HashSet<RecordField> hs = new HashSet<RecordField>(writeSets.get(it));
        if(hs.retainAll(readSet)) {
            return false;
        }
    }

    // write set of transaction in active Set interacts read set or write set
    for(Transaction activeTx: active) {
        HashSet<RecordField> hs1 = new HashSet<RecordField>(activeTx.getWriteSet());
        HashSet<RecordField> hs2 = new HashSet<RecordField>(hs1);
        if(hs1.retainAll(readSet)) {
            return false;
        }
        if(hs2.retainAll(writeSet)) {
            return false;
        }
    }

    return true;
}
```

make copies of tnc and active tx

# validate()

TransactionMgr

```
public boolean validate(Transaction tx) {
    long start = tx.getStartTn();
    long finish = 0;
    HashSet<Transaction> active;
    synchronized (tncLock) {
        synchronized (activeLock) {
            finish = tnc;
            active = new HashSet<Transaction>(activeSets);
            activeSets.add(tx);
        }
    }
    HashSet<RecordField> readSet = tx.getReadSet();
    HashSet<RecordField> writeSet = tx.getWriteSet();

    // write set of transaction from start tn+1 to finish tn intersects read set
    for(long it=start+1; it<=finish; ++it) {
        HashSet<RecordField> hs = new HashSet<RecordField>(writeSets.get(it));
        if(hs.retainAll(readSet)) {
            return false;
        }
    }

    // write set of transaction in active Set intersects read set or write set
    for(Transaction activeTx: active) {
        HashSet<RecordField> hs1 = new HashSet<RecordField>(activeTx.getWriteSet());
        HashSet<RecordField> hs2 = new HashSet<RecordField>(hs1);
        if(hs1.retainAll(readSet)) {
            return false;
        }
        if(hs2.retainAll(writeSet)) {
            return false;
        }
    }

    return true;
}
```

**follow condition (2)**

(2) The write set of  $T_i$  does not intersect the read set of  $T_j$ , and  $T_i$  completes its write phase before  $T_j$  starts its write phase.

# validate()

TransactionMgr

```
public boolean validate(Transaction tx) {
    long start = tx.getStartTn();
    long finish = 0;
    HashSet<Transaction> active;
    synchronized (tncLock) {
        synchronized (activeLock) {
            finish = tnc;
            active = new HashSet<Transaction>(activeSets);
            activeSets.add(tx);
        }
    }
    HashSet<RecordField> readSet = tx.getReadSet();
    HashSet<RecordField> writeSet = tx.getWriteSet();

    // write set of transaction from start tn+1 to finish tn interacts read set
    for(long it=start+1; it<=finish; ++it) {
        HashSet<RecordField> hs = new HashSet<RecordField>(writeSets.get(it));
        if(hs.retainAll(readSet)) {
            return false;
        }
    }

    // write set of transaction in active Set interacts read set or write set
    for(Transaction activeTx: active) {
        HashSet<RecordField> hs1 = new HashSet<RecordField>(activeTx.getWriteSet());
        HashSet<RecordField> hs2 = new HashSet<RecordField>(hs1);
        if(hs1.retainAll(readSet)) {
            return false;
        }
        if(hs2.retainAll(writeSet)) {
            return false;
        }
    }

    return true;
}
```

**follow condition (3)**

(3) The write set of  $T_i$  does not intersect the read set or the write set of  $T_j$ , and  $T_i$  completes its read phase before  $T_j$  completes its read phase.



# end validation

```
tend = (  
  {finish tn := tnc;  
   finish active := (make a copy of active);  
   active := active  $\cup$  {id of this transaction}};  
   valid := true;  
for t from start tn + 1 to finish tn do  
  if (write set of transaction with transaction number t intersects read set)  
    then valid := false;  
for i  $\in$  finish active do  
  if (write set of transaction  $T_i$  intersects read set or write set)  
    then valid := false;  
if valid  
  then (  
    (write phase);  
    {tnc := tnc + 1;  
     tn := tnc;  
     active := active — {id of this transaction}};  
    (cleanup))  
  else (  
    {active := active — {id of transaction}};  
    {backup}}).
```

## Transaction

```
TransactionMgr  
public void finishedCommit(Set<RecordField> writeSet) {  
  synchronized(tncLock) {  
    tnc++;  
    writeSets.put(tnc, new HashSet<RecordField>(writeSet));  
  }  
}  
  
public void commit() {  
  if(VanillaDb.txMgr().validate(this)) {  
    commitCertify();  
    commitWorkspace();  
    for (TransactionLifecycleListener l : lifecycleListeners)  
      l.onTxCommit(this);  
  
    if (logger.isLoggable(Level.FINE))  
      logger.fine("transaction " + txNum + " committed");  
  
    VanillaDb.txMgr().finishedCommit(writeSet.keySet());  
    VanillaDb.txMgr().endCommit(this);  
  } else {  
    VanillaDb.txMgr().endCommit(this);  
    throw new ValidAbortException("abort tx." + this.txNum + " for no validation");  
  }  
}
```

add to txMgr write sets

# end validation

```
tend = (  
  {finish tn := tnc;  
   finish active := (make a copy of active);  
   active := active  $\cup$  {id of this transaction}};  
   valid := true;  
for t from start tn + 1 to finish tn do  
  if (write set of transaction with transaction number t intersects read set)  
  then valid := false;  
for i  $\in$  finish active do  
  if (write set of transaction  $T_i$  intersects read set or write set)  
  then valid := false;  
if valid  
then (  
  (write phase);  
  {tnc := tnc + 1;  
   tn := tnc;  
   active := active — {id of this transaction}};  
  (cleanup))  
else (  
  {active := active — {id of this transaction}};  
  {backup}}).
```

## TransactionMgr

```
public void endCommit(Transaction tx) {  
  synchronized(activeLock) {  
    activeSets.remove(tx);  
  }  
}
```

## Transaction

```
public void commit() {  
  if(VanillaDb.txMgr().validate(this)) {  
    commitCertify();  
    commitWorkspace();  
    for (TransactionLifecycleListener l : lifecycleListeners)  
      l.onTxCommit(this);  
  
    if (logger.isLoggable(Level.FINE))  
      logger.fine("transaction " + txNum + " committed");  
  
    VanillaDb.txMgr().finishedCommit(writeSet.keySet());  
    VanillaDb.txMgr().endCommit(this);  
  } else {  
    VanillaDb.txMgr().endCommit(this);  
    throw new ValidAbortException("abort tx." + this.txNum + " for no validation");  
  }  
  remove itself from active set  
}
```

# end validation

```
tend = (  
  {finish tn := tnc;  
   finish active := (make a copy of active);  
   active := active  $\cup$  {id of this transaction}};  
   valid := true;  
for t from start tn + 1 to finish tn do  
  if (write set of transaction with transaction number t intersects read set)  
  then valid := false;  
for i  $\in$  finish active do  
  if (write set of transaction  $T_i$  intersects read set or write set)  
  then valid := false;  
if valid  
then (  
  (write phase);  
  {tnc := tnc + 1;  
   tn := tnc;  
   active := active — {id of this transaction}};  
  (cleanup))  
else (  
  {active := active — {id of transaction}};  
  {backup}}).
```

## Transaction

```
public void commit() {  
  if(VanillaDb.txMgr().validate(this)) {  
    commitCertify();  
    commitWorkspace();  
    for (TransactionLifecycleListener l : lifecycleListeners)  
      l.onTxCommit(this);  
  
    if (logger.isLoggable(Level.FINE))  
      logger.fine("transaction " + txNum + " committed");  
  
    VanillaDb.txMgr().finishedCommit(writeSet.keySet());  
    VanillaDb.txMgr().endCommit(this);  
  } else {  
    VanillaDb.txMgr().endCommit(this);  
    throw new ValidAbortException("abort tx." + this.txNum + " for no validation");  
  }  
}
```

no validation throw an error and will be rollback

# write phase

**for**  $n \in \text{write set}$  **do**  $\text{exchange}(n, \text{copies}[n])$ .  
**for**  $n \in \text{delete set}$  **do**  $\text{delete}(n)$ ;

## Transaction

```
public void commit() {  
    if(VanillaDb.txMgr().validate(this)) {  
        commitCertify();  
        commitWorkspace();  
        for (TransactionLifecycleListener l : lifecycleListeners)  
            l.onTxCommit(this);  
  
        if (logger.isLoggable(Level.FINE))  
            logger.fine("transaction " + txNum + " committed");  
  
        VanillaDb.txMgr().finishedCommit(writeSet.keySet());  
        VanillaDb.txMgr().endCommit(this);  
    } else {  
        VanillaDb.txMgr().endCommit(this);  
        throw new ValidAbortException("abort tx." + this.txNum + " for no validation");  
    }  
}
```

```
public void commitCertify() {  
    this.cCertified = true;  
}
```

# write phase

**for**  $n \in \text{write set}$  **do**  $\text{exchange}(n, \text{copies}[n])$ .  
**for**  $n \in \text{delete set}$  **do**  $\text{delete}(n)$ ;

## Transaction

```
public void commit() {  
  
    if(VanillaDb.txMgr().validate(this)) {  
        commitCertify();  
        commitWorkspace();  
        for (TransactionLifecycleListener l : lifecycleListeners)  
            l.onTxCommit(this);  
  
        if (logger.isLoggable(Level.FINE))  
            logger.fine("transaction " + txNum + " committed");  
  
        VanillaDb.txMgr().finishedCommit(writeSet.keySet());  
        VanillaDb.txMgr().endCommit(this);  
    } else {  
        VanillaDb.txMgr().endCommit(this);  
        throw new ValidAbortException("abort tx." + this.txNum + " for no validation");  
    }  
}
```



# write phase

**for**  $n \in \text{write set}$  **do**  $\text{exchange}(n, \text{copies}[n])$ .  
**for**  $n \in \text{delete set}$  **do**  $\text{delete}(n)$ ;

## Transaction

```
private void commitWorkspace() {  
    for (Map.Entry<RecordField, Constant> entry: writeSet.entrySet()) {  
        RecordField rfield = entry.getKey();  
        Constant val = entry.getValue();  
        TableInfo ti = VanillaDb.catalogMgr().getTableInfo(rfield.tblName, this);  
        RecordFile rfile = ti.open(this, true);  
        rfile.moveToRecordId(rfield.rid);  
        rfile.setVal(rfield.fldName, val);  
        rfile.close();  
    }  
  
    for (RecordField rfield: deleteSet) {  
        TableInfo ti = VanillaDb.catalogMgr().getTableInfo(rfield.tblName, this);  
        RecordFile rfile = ti.open(this, true);  
        rfile.delete(rfield.rid);  
        rfile.close();  
    }  
}
```

**get a RecordFile of the table to be modified**

# write phase

**for**  $n \in \text{write set}$  **do**  $\text{exchange}(n, \text{copies}[n])$ .  
**for**  $n \in \text{delete set}$  **do**  $\text{delete}(n)$ ;

## Transaction

```
private void commitWorkspace() {  
    for (Map.Entry<RecordField, Constant> entry: writeSet.entrySet()) {  
        RecordField rfield = entry.getKey();  
        Constant val = entry.getValue();  
        TableInfo ti = VanillaDb.catalogMgr().getTableInfo(rfield.tblName, this);  
        RecordFile rfile = ti.open(this, true);  
        rfile.moveToRecordId(rfield.rid);  
        rfile.setVal(rfield.fldName, val);  
        rfile.close();  
    }  
  
    for(RecordField rfield: deleteSet) {  
        TableInfo ti = VanillaDb.catalogMgr().getTableInfo(rfield.tblName, this);  
        RecordFile rfile = ti.open(this, true);  
        rfile.delete(rfield.rid);  
        rfile.close();  
    }  
}
```

modify through the original path

# write phase

**for**  $n \in \text{write set}$  **do**  $\text{exchange}(n, \text{copies}[n])$ .  
**for**  $n \in \text{delete set}$  **do**  $\text{delete}(n)$ ;

## Transaction

```
private void commitWorkspace() {
    for (Map.Entry<RecordField, Constant> entry: writeSet.entrySet()) {
        RecordField rfield = entry.getKey();
        Constant val = entry.getValue();
        TableInfo ti = VanillaDb.catalogMgr().getTableInfo(rfield.tblName, this);
        RecordFile rfile = ti.open(this, true);
        rfile.moveToRecordId(rfield.rid);
        rfile.setVal(rfield.fldName, val);
        rfile.close();
    }

    for(RecordField rfield: deleteSet) {
        TableInfo ti = VanillaDb.catalogMgr().getTableInfo(rfield.tblName, this);
        RecordFile rfile = ti.open(this, true);
        rfile.delete(rfield.rid);
        rfile.close();
    }
}

public void setVal(String fldName, Constant val) {
    if (tx.isReadOnly() && !isTempTable())
        throw new UnsupportedOperationException();
    Type fldType = ti.schema().type(fldName);

    Constant v = val.castTo(fldType);
    if (Page.size(v) > Page.maxSize(fldType))
        throw new SchemaIncompatibleException();
    if (!tx.rollbackCertified() && !tx.commitCertified() && !isTempTable()) {
        tx.putVal(ti.tableName(), currentRecordId(), fldName, v);
    } else {
        rp.setVal(fldName, v);
    }
}
```

## RecordFile



# write phase

**for**  $n \in \text{write set}$  **do**  $\text{exchange}(n, \text{copies}[n])$ .  
**for**  $n \in \text{delete set}$  **do**  $\text{delete}(n)$ ;

## Transaction

```
private void commitWorkspace() {  
    for (Map.Entry<RecordField, Constant> entry: writeSet.entrySet()) {  
        RecordField rfield = entry.getKey();  
        Constant val = entry.getValue();  
        TableInfo ti = VanillaDb.catalogMgr().getTableInfo(rfield.tblName, this);  
        RecordFile rfile = ti.open(this, true);  
        rfile.moveToRecordId(rfield.rid);  
        rfile.setVal(rfield.fldName, val);  
        rfile.close();  
    }  
  
    for (RecordField rfield: deleteSet) {  
        TableInfo ti = VanillaDb.catalogMgr().getTableInfo(rfield.tblName, this);  
        RecordFile rfile = ti.open(this, true);  
        rfile.delete(rfield.rid);  
        rfile.close();  
    }  
}
```

**modify through the original path**

# write phase

**for**  $n \in \text{write set}$  **do**  $\text{exchange}(n, \text{copies}[n])$ .  
**for**  $n \in \text{delete set}$  **do**  $\text{delete}(n)$ ;

## Transaction

```
private void commitWorkspace() {
    for (Map.Entry<RecordId, RecordField> en : recordFields.entrySet()) {
        Constant val = en.getValue();
        TableInfo ti = Val.getTableInfo(val);
        RecordFile rfile = new RecordFile(ti, val);
        rfile.moveToRecordFile(ti.getTableFile());
        rfile.setVal(rfile.getVal());
        rfile.close();
    }

    for (RecordField rfield : recordFields.values()) {
        TableInfo ti = rfield.getTableInfo();
        RecordFile rfile = new RecordFile(ti, rfield.getVal());
        rfile.delete(rfield.getVal());
        rfile.close();
    }
}
```

```
public void delete() {
    if (tx.isReadOnly() && !isTempTable())
        throw new UnsupportedOperationException();

    if (!tx.rollbackCertified() && !tx.commitCertified()) {
        tx.delete(ti.tableName(), currentRecordId());
    } else {
        if (fhp == null)
            fhp = openHeaderForModification();

        // Log that this logical operation starts
        RecordId deletedRid = currentRecordId();
        tx.recoveryMgr().logLogicalStart();

        // Delete the current record
        rp.delete(fhp.getLastDeletedSlot());
        fhp.setLastDeletedSlot(currentRecordId());

        // Log that this logical operation ends
        tx.recoveryMgr().logRecordFileDeletionEnd(ti.tableName(), deletedRid.block().number(), deletedRid.id());

        // Close the header (release the header lock)
        closeHeader();
    }
}
```

## RecordFile

# write phase

**for**  $n \in \text{write set}$  **do**  $\text{exchange}(n, \text{copies}[n])$ .  
**for**  $n \in \text{delete set}$  **do**  $\text{delete}(n)$ ;

## Transaction

```
public void rollback() {  
    rollbackCertify();  
  
    for (TransactionLifecycleListener l : lifecycleListeners) {  
        l.onTxRollback(this);  
    }  
  
    if (logger.isLoggable(Level.FINE))  
        logger.fine("transaction " + txNum + " rolled back");  
}
```

```
public void rollbackCertify() {  
    this.rCertified = true;  
}
```

# Outline

- Paper introduction
  - Motivation and Problem
  - Main Idea
  - Conclusion
- Implementation
- **Evaluation and Experiments**
- Conclusion

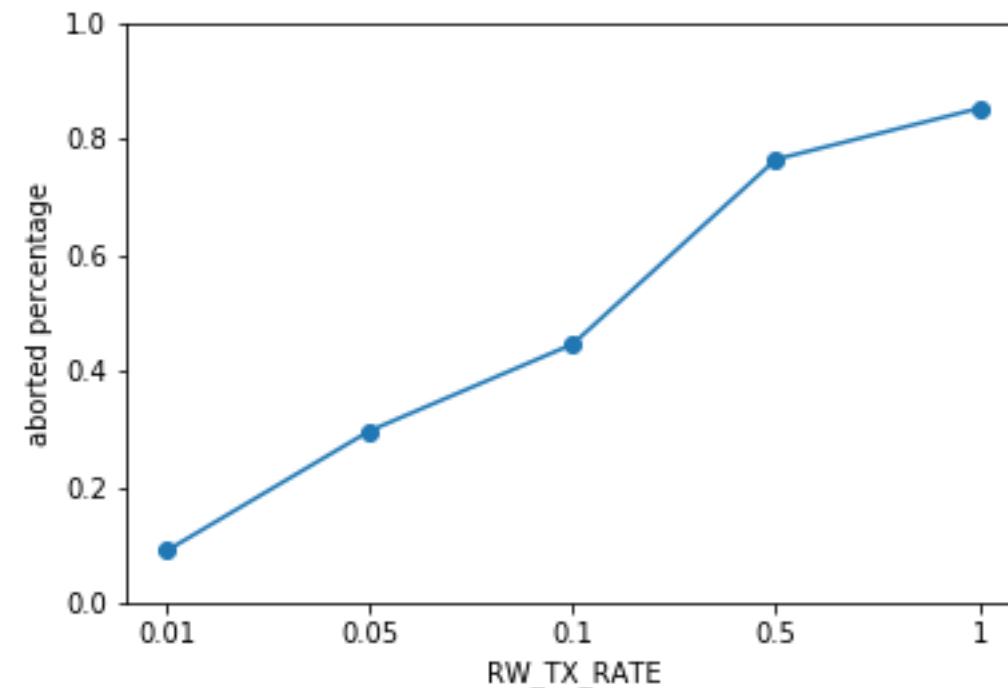

# Evaluation and Experiment

- environment



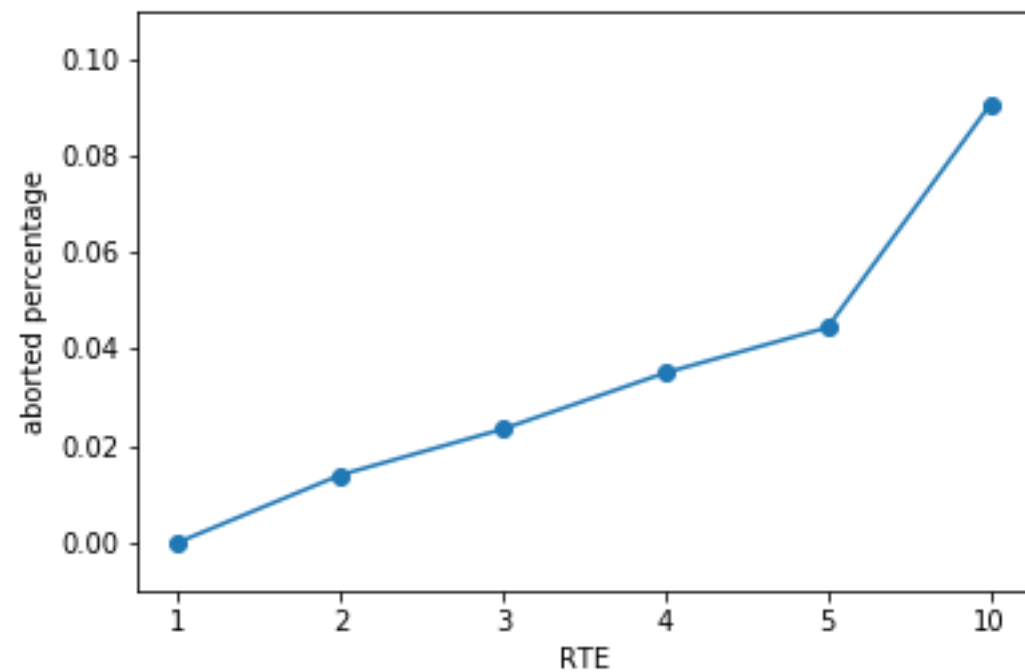
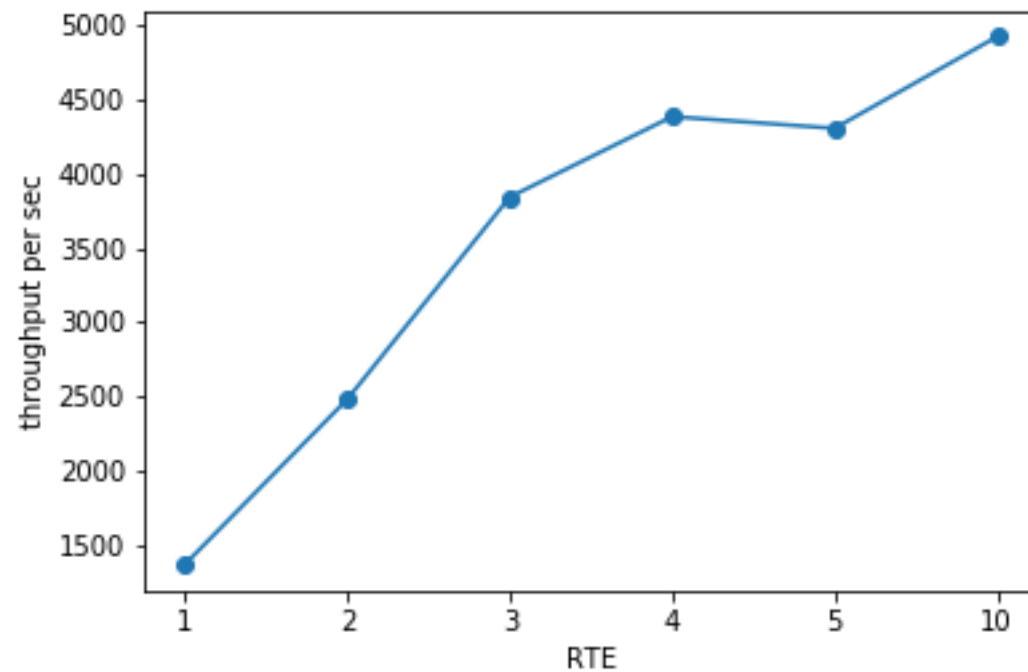
# Evaluation and Experiment

- Exp. I - Change RW\_TX\_RATE to verify



# Evaluation and Experiment

- Exp. II - Change RTE to verify



# Outline

- Paper introduction
  - Motivation and Problem
  - Main Idea
  - Conclusion
- Implementation
- Evaluation and Experiments
- Conclusion



# Conclusion

- the paper did not mention the detail, so we spend more time on thinking what will happen.
- our OCC experiment shows that the curve is fit the OCC property.